

Installing Kubernetes Using Kubeadm

```
=====
=
#
Topic: Setting up Control-plane/Master node
#####
1)
# Install container runtime - Docker
Follow this source.
Source:
https://docs.docker.com/install/linux/docker-ce/ubuntu/

2)
# Install Kubeadm
Follow this source or below commands.

Source:
https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/

Commands:
- sudo apt-get update

- sudo apt-get install -y apt-transport-https curl

- Add repository keys to download tools:
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg |
sudo apt-key add -

- Add repository:
cat <<EOF | sudo tee /etc/apt/sources.list.d/kubernetes.list
deb https://apt.kubernetes.io/ kubernetes-xenial main
EOF
- sudo apt-get update
```

- `sudo apt-get install -y kubelet kubeadm kubectl`
- `sudo apt-mark hold kubelet kubeadm kubectl`

#3) Initialize a new Kubernetes cluster

`$ sudo kubeadm init`

NOTE: Save this output as it's required to add worker nodes.

4)

As per the instruction from 'kubeadm init' command output,
To make kubectl work for your non-root user, run these
commands.

`mkdir -p $HOME/.kube`

`sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config`

`sudo chown $(id -u):$(id -g) $HOME/.kube/config`

#5) Verify if cluster is initialized succussfully

`$ kubectl get nodes`

O/P:

NAME	STATUS	ROLES	AGE	VERSION
node1	NotReady	master	2m43s	v1.12.1

#9) Run the following kubectl command to find the reason why the cluster STATUS is showing as NotReady.

- This command shows all Pods in all namespaces - this includes system Pods in the system (kube-system) namespace.

- As we can see, none of the coredns Pods are running

- This is preventing the cluster from entering the Ready state, and is happening because we haven't created the Pod network yet.

O/P:

`$ kubectl get pods --all-namespaces`

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-...vt	0/1	ContainerCreating	0	8m33s
kube-system	coredns-...xw	0/1	ContainerCreating	0	8m33s
kube-system	etcd...	1/1	Running	0	7m46s
kube-system	kube-api...	1/1	Running	0	7m36s

#7) Create Pod Network. You must install a pod network add-on so that your pods can communicate with each other. (As per kubeadm init output)

Source:

<https://www.weave.works/docs/net/latest/kubernetes/kube-addon/#install>

```
$ kubectl apply -f "https://cloud.weave.works/k8s/net?k8s-version=$(kubectl version | base64 | tr -d '\n')"
```

#8) Check if the status of Master is changed from 'NotReady' to 'Ready'

```
$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
node1	Ready	master	3m51s	v1.12.1

GREAT - the cluster is ready and all dns system pods are now working. Cluster is ready now.

Now that the cluster is up and running, it's time to add some worker-nodes.

#

Topic: Worker Node Setup & Joining to the cluster:

#####

1

Create a worker node machine in GCP / AWS cloud platform.

2

Install kubeadm

<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>

3

Install container runtime

<https://docs.docker.com/install/linux/docker-ce/ubuntu/>

4

To bootstrap a Kubernetes worker node and join it to the cluster run below command from \$kubeadm init output.

Note: bootstrap: It automatically installs kubectl and kubelet.

Note: Below Bootstrap token will be different for your Control-plane/Master node. Use the one which you have copied at step:3.

```
kubeadm join 10.128.0.18:6443 --token
9ril81.t4k4sqh1ionqv1om \
--discovery-token-ca-cert-hash
sha256:de57d9e08877db501a8b503db3ee91596f8f5657878
c3087bc0343ece7df3eb2
```

Verify node Join (Run below in Control-plane node)

\$ kubectl get nodes

NAME	STATUS	ROLES	AGE	VERSION
control-plane	Ready	master	26m	v1.16.3
worker-node1	Ready	<none>	3m18s	v1.16.3

\$ kubectl get nodes -o wide

--> this will display IP, OS, Kernel and more details about all Nodes

#

Project-1 [Nginx]

#####

Deploying/Creating a pod

#####

#

1.) Create Pod manifest file

\$ mkdir nginx

\$ vim pod.yaml

pod.yaml

=====

apiVersion: v1

kind: Pod

metadata:

name: nginx-pod

labels:

env: prod

version: v1.2.3

spec:

containers:

- name: nginx-container

image: nginx

ports:

- containerPort: 80

pod.yaml - Manifest file description:

- Straight away we can see four top-level resources.

- .apiVersion
- .kind
- .metadata
- .spec

--> .apiVersion:

- Tells API Server about what version of Yaml is used to create the object (Pod object in this case)
- Pods are currently in v1 API group

--> .kind:

- Tells us the kind of object being deployed. In this case we are creating POD object.
- It tells control plane what type of object is being defined.

--> .metadata:

- this section again has two sub-sections i.e name & labels
- You can name the Pod using "name" key.
- Using labels, we can identify a particular pod.

--> .spec:

- This is where we specify details about the containers that will run in the Pod.
- In this section we specify container name, image, ports ..etc.

#

2.) Creating a Pod

- Check if all Nodes are ready before creating a Pod

```
$ kubectl get nodes
```

- This POSTs the manifest file to API server and deploy/create a Pod from it

```
$ kubectl apply -f pod.yml
```

Note: Your Pod has been scheduled to a healthy node in the cluster and

is being monitored by the local kubelet process on the node.

Introspecting Running Pods

- Get IP and worker node of the Pod

```
$ kubectl get pod -o wide
```

- Launch nginx server application running in the Pod from Controle-plane node

```
$ curl http://10.44.0.1:80
```

```
$ curl http://POD-IP:Server-Port
```

- You can also login into the Pod container to get more information.

```
$ kubectl exec -it nginx-pod /bin/bash
```

Note: Let's add some code and launch our nginx application

```
- $ echo "Gamut Gurus Technologies" > /usr/share/nginx/html/index.html
```

- Launch nginx application

```
$ curl http://10.44.0.1:80
```

- Login into a specific container in case you have multi container Pod

using --container or -c option.

```
$ kubectl exec -it nginx-pod --container nginx-container /bin/bash
```

#

3.) Deleting a Pod

```
$ kubectl get pods
```

```
$ kubectl delete pods nginx-pod
```

```
$ kubectl delete -f pod.yml
```

NOTE:

kubelet takes the PodSpec and is responsible for pulling all images and starting all containers in the Pod.

What Next?

- If a Pod fails, it is not automatically rescheduled. Because of this, we usually deploy them via higher-level object such as Deployments.

- This adds things like "scalability" (scale-up/down), "self-healing", "rolling updates" and "roll backs" and makes Kubernetes so powerful.

Misc. CMDs:

- Get full copy of the Pod manifest from cluster store. desired state is (.spec) and observed state will be under (.status)

```
$ kubectl get pod -o yaml
```

- Check if Pod is created

```
$ kubectl get pods
```

```
$ kubectl get pods --watch (monitor the status continuously)
```

- Another great Kubernetes introspection command. Provides Pods(object's) lifecycle events.

```
$ kubectl describe pod nginx-pod
```

--wk-end-2PM-batch--

#

Project-2 [Nginx]

```
#####
```

Creating Deployments & Services

```
#####
```

- Pods don't self-heal, they don't scale, and they don't allow for easy updates

- Deployments do all things like
 - "auto scale" (scale-up/down)
 - "self-heal"
 - "rolling updates"
 - "roll backs"

- That's why we almost always deploy Pods via 'Deployments'


```
#
Creating Deployments
-----
# List all nodes in K8s cluster
$ kubectl get nodes

# List all pods in K8s cluster
$ kubectl get pods

# Create the deployment
$ kubectl create -f deploy-nginx.yml
```

```
vim deploy-nginx.yml
-----
apiVersion: apps/v1
kind: Deployment
metadata:
  name: flipkart-prod-deploy
  labels:
    app: flipkart-prod
spec:
  replicas: 6
  selector:
    matchLabels:
      app: flipkart-pod-template
  template:
    metadata:
      labels:
        app: flipkart-pod-template
    spec:
      containers:
        - name: nginx-container
          image: nginx
          ports:
            - containerPort: 80
```

Creating deployment

```
$ kubectl create -f deploy-nginx.yml
```

Check pod creations

```
$ kubectl get pods --watch
```

Login to pods and verify nginx application

```
$ kubectl get pods -o wide
```

```
$ kubectl exec -it nginx-deploy-5f654bcccd-27xtg /bin/bash
```

launch application from individual Pod

```
$ curl http://10.44.0.1:80
```

```
$ curl http://pod_ip:80
```

Testing Self-healing capability

If you delete some Pods, Kubernetes can automatically re-create the same for us to make sure given no. of Pods are always running.

- Delete the Pods

```
$ kubectl delete pods POD_NAME1 POD_NAME2
```

- Check if the Pods are re-created

```
$ kubectl get pods
```

Creating service to expose the application to outside world and setting up load balancer

=====

==

```
$ vim service-nginx.yml
```

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
  name: flipkart-service
  labels:
    app: flipkart-pod-service
```

```
spec:
  selector:
    app: flipkart-pod-template
  type: NodePort
  ports:
    - nodePort: 31000
      port: 80
      targetPort: 80
```

Create the service

```
$ kubectl create -f service-nginx.yml
```

Enable networking

Click on Navigation menu(three lines on top left) --> Go to VPC Network --> Firewall rules --> select on one existing rule --> edit --> Source IP ranges --> 0.0.0.0/0 --> In 'Specified protocols and ports', write this range "0-65535"

Access the application from browser using worker-node port

<http://34.93.139.52:31000/>

<http://WorkerNodeIP:Node-port/>

Project-3 [GamutKart]

=====

Creating deployment for GamutKart

```
$ vim deploy-gamutkart.yml
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: gamutkart-deploy
```

```
  labels:
    app: gamutkart-app
spec:
  replicas: 8
  selector:
    matchLabels:
      app: gamutkart-app
  template:
    metadata:
      labels:
        app: gamutkart-app
    spec:
      containers:
        - name: gamutkart-container
          image: nageshvkn/gamutkart-img
          ports:
            - containerPort: 8080
          command: ["/bin/sh"]
          args: ["-c", "/root/apache-tomcat-8.5.38/bin/startup.sh;
while true; do sleep 1; done;"]
```

```
# Execute deployment
$ kubectl create -f deploy-gamutkart.yml
```

```
# Creating service for GamutKart
$ vim service-gamutkart.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: gamutkart-service
  labels:
    app: gamutkart-app
spec:
  selector:
    app: gamutkart-app
  type: NodePort
```

```
ports:
- nodePort: 31000
  port: 8080
  targetPort: 8080
```

```
# Creating the service
$ kubectl create -f service-gamutkart.yml
```

```
#
# Enable networking
TODO:
Go to VPC Network --> Firewall --> select on one existing rule
--> edit --> Source IP ranges
--> 0.0.0.0/0 --> In "Specified protocols and ports", write this
range "0-65535"
```

```
Note:
Kubernetes Port Range: 30,000 - 32,767
```

```
#
3.) Deleting a Pod
$ kubectl get pods
$ kubectl delete pods nginx-pod
$ kubectl delete -f pod.yml
```

```
# Misc:
4.) Get all nodes IPs in Kubernetes cluster
$ kubectl get nodes -o wide
```

```
#
List Deployments & Service
$ kubectl get deployment
$ kubectl get svc (Or service)
```

#

5.) Deleting Deployment & Service

```
$ kubectl delete -f deploy-gamutkart.yaml(deployment yaml  
file name)
```

```
$ kubectl delete -f service-gamutkart.yml( service yaml file  
name)
```

```
$ kubectl delete deployment <deployment-name>
```

```
$ kubectl delete service <service-name>
```

#Scaleup Pods

```
$ kubectl scale deployments/gamutkart-deploy --replicas=2
```

8:30 PM daily

--

- Creatng a POD on a particular WN
- Autoscale example
- Assign different names to different PODs manually
- Create different replicas for different PODs
- select a deployment as part ofa serice
- How to list all the API versions w.r.t Object implementation.
-